

Introduction to SQL Queries

Pair Programming Exercise for DSE5002

Peter Gyorda, April 12 2025

We will connect to the chinook database and work through some queries

These are queries made to a single table, we will see later how to join tables together and create subqueries

The commands used in queries are:

SELECT --specify the variables or columns required

FROM--specify the table to obtain data from

LIMIT-- restrict the number of lines returned to a desired total N

WHERE-- this is a filtering function carried out on row elements, we can use AND, OR and NOT within the Where

ORDER BY-- This is a sorting function, it can sort ascending or descending, and we can sort on multiple variables

GROUP BY--this is a grouping function

HAVING--Having is a filtering operation on group members

MAX(), MIN(), AVG(), SUM(), COUNT() are aggregating functions used with GROUP BY,

Source material

"Learning SQL", Beaulieu, O'Reilly 2005

<https://www.sqlitetutorial.net/> - explains queries using the chinook database, albeit in the SQLite database system. The SELECT system used for queries is pretty standard for most SQL databases, the other aspects and commands seem to be a bit more variable from one server program to another.

That said, there are minor differences in variable names between the chinook database in postgres and the tutorial for SQLite, watch for underscores and pluralization (track vs tracks, etc). I have fixed all the examples shown here.

Connect to the Chinook Database

and figure out what we have in it

Set up the required libraries

```
In [57]: # First, install the required PostgreSQL adapter
!pip install psycopg2-binary

#import psycopg2
import sqlalchemy

# we will want Pandas for the data frame structure

import pandas as pd
```

Requirement already satisfied: psycopg2-binary in c:\users\petea\anaconda3\lib\site-packages (2.9.10)

```
In [59]: # Alter this to reflect your username and password,    this is for postgres on the same machine

engine=sqlalchemy.create_engine('postgresql://todd:password@localhost:5432/chinook')
```

what tables do we have

We can use a SELECT command to look for the table_name values in a built-in database called information_schema, in a table called tables.

This is a database and table that are built into postgres to hold information. It holds a lot of info, but our table names are all in the first 15 lines

```
In [61]: pd.read_sql_query("SELECT table_name FROM information_schema.tables LIMIT 15",engine)
```

```
Out[61]:
```

	table_name
0	artist
1	album
2	employee
3	customer
4	invoice
5	invoice_line
6	track
7	playlist
8	playlist_track
9	genre
10	pg_statistic
11	pg_type
12	media_type
13	pg_foreign_table
14	pg_authid

```
In [65]: # Looking at the customer table, but only first 5 rows
```

```
pd.read_sql_query("SELECT * FROM customer LIMIT 5",engine)
```

Out[65]:

	customer_id	first_name	last_name	company	address	city	state	country	postal_code	phone	fax
0	1	Luís	Gonçalves	Embraer - Empresa Brasileira de Aeronáutica S.A.	Av. Brigadeiro Faria Lima, 2170	São José dos Campos	SP	Brazil	12227-000	+55 (12) 3923-5555	+55 (12) 3923-5566
1	2	Leonie	Köhler	None	Theodor-Heuss-Straße 34	Stuttgart	None	Germany	70174	+49 0711 2842222	None lec
2	3	François	Tremblay	None	1498 rue Bélanger	Montréal	QC	Canada	H2G 1A7	+1 (514) 721-4711	None ft
3	4	Bjørn	Hansen	None	Ullevålsveien 14	Oslo	None	Norway	0171	+47 22 44 22 22	None bjor
4	5	František	Wichterlová	JetBrains s.r.o.	Klanova 9/506	Prague	None	Czech Republic	14700	+420 2 4172 5555	+420 2 4172 5555

In [67]: *#restrict this to only customer_id, first and last names*

```
pd.read_sql_query("SELECT customer_id, first_name, last_name FROM customer LIMIT 8", engine)
```

Out[67]:

	customer_id	first_name	last_name
0	1	Luís	Gonçalves
1	2	Leonie	Köhler
2	3	François	Tremblay
3	4	Bjørn	Hansen
4	5	František	Wichterlová
5	6	Helena	Holý
6	7	Astrid	Gruber
7	8	Daan	Peeters

QUESTION/ACTION

Figure out what the table "invoices" looks like, display the first 5 lines of it so you can see the content

In [69]: `pd.read_sql_query("SELECT * FROM invoice LIMIT 5", engine)`

Out[69]:

	invoice_id	customer_id	invoice_date	billing_address	billing_city	billing_state	billing_country	billing_postal_code	total
0	1	2	2021-01-01	Theodor-Heuss-Straße 34	Stuttgart	None	Germany	70174	1.98
1	2	4	2021-01-02	Ullevålsveien 14	Oslo	None	Norway	0171	3.96
2	3	8	2021-01-03	Grétrystraat 63	Brussels	None	Belgium	1000	5.94
3	4	14	2021-01-06	8210 111 ST NW	Edmonton	AB	Canada	T6G 2C7	8.91
4	5	23	2021-01-11	69 Salem Street	Boston	MA	USA	2113	13.86

Question/Action

Show the variables customer_id, billing_country and total for the first 12 lines of invoice

```
In [71]: pd.read_sql_query("SELECT customer_id, billing_country, total FROM invoice LIMIT 12", engine)
```

```
Out[71]:    customer_id  billing_country  total
```

	customer_id	billing_country	total
0	2	Germany	1.98
1	4	Norway	3.96
2	8	Belgium	5.94
3	14	Canada	8.91
4	23	USA	13.86
5	37	Germany	0.99
6	38	Germany	1.98
7	40	France	1.98
8	42	France	3.96
9	46	Ireland	5.94
10	52	United Kingdom	8.91
11	2	Germany	13.86

Ordering or Sorting Results

```
In [73]: pd.read_sql_query("SELECT * FROM track ORDER BY Milliseconds LIMIT 12", engine)
```

Out[73]:

	track_id	name	album_id	media_type_id	genre_id	composer	milliseconds	bytes	unit_price
0	2461	É Uma Partida De Futebol	200	1	1	Samuel Rosa	1071	38747	0.99
1	168	Now Sports	18	1	4	None	4884	161266	0.99
2	170	A Statistic	18	1	4	None	6373	211997	0.99
3	178	Oprah	18	1	4	None	6635	224313	0.99
4	3304	Commercial 1	258	1	17	L. Muggerud	7941	319888	0.99
5	172	The Real Problem	18	1	4	None	11650	387360	0.99
6	3310	Commercial 2	258	1	17	L. Muggerud	21211	850698	0.99
7	2241	Bossa	184	1	17	None	29048	967098	0.99
8	1086	Casinha Feliz	85	1	10	Gilberto Gil	32287	1039615	0.99
9	246	Mateus Enter	24	1	7	Chico Science	33149	1103013	0.99
10	975	Deixa Entrar	78	1	7	None	33619	1095012	0.99
11	2797	Homem Primata (Vinheta)	224	1	4	Titãs	34168	1124909	0.99

In [11]:

```
# reversed order sort

# add DESC to sort descending, ASC to sort ascending

pd.read_sql_query("SELECT * FROM track ORDER BY Milliseconds DESC LIMIT 12",engine)
```

Out[11]:

	track_id	name	album_id	media_type_id	genre_id	composer	milliseconds	bytes	unit_price
0	2820	Occupation / Precipice	227	3	19	None	5286953	1054423946	1.99
1	3224	Through a Looking Glass	229	3	21	None	5088838	1059546140	1.99
2	3244	Greetings from Earth, Pt. 1	253	3	20	None	2960293	536824558	1.99
3	3242	The Man With Nine Lives	253	3	20	None	2956998	577829804	1.99
4	3227	Battlestar Galactica, Pt. 2	253	3	20	None	2956081	521387924	1.99
5	3226	Battlestar Galactica, Pt. 1	253	3	20	None	2952702	541359437	1.99
6	3243	Murder On the Rising Star	253	3	20	None	2935894	551759986	1.99
7	3228	Battlestar Galactica, Pt. 3	253	3	20	None	2927802	554509033	1.99
8	3248	Take the Celestra	253	3	20	None	2927677	512381289	1.99
9	3239	Fire In Space	253	3	20	None	2926593	536784757	1.99
10	3232	The Long Patrol	253	3	20	None	2925008	513122217	1.99
11	3235	The Magnificent Warriors	253	3	20	None	2924716	570152232	1.99

In [75]:

```
# sort by two variables
pd.read_sql_query("SELECT * FROM track ORDER BY composer ASC, milliseconds DESC LIMIT 12",engine)
```

Out[75]:

	track_id	name	album_id	media_type_id	genre_id	composer	milliseconds	bytes	unit_price
0	2108	Children Of The Grave	174	1	3	A. F. Iommi, W. Ward, T. Butler, J. Osbourne	357067	11626740	0.99
1	2109	Paranoid	174	1	3	A. F. Iommi, W. Ward, T. Butler, J. Osbourne	176352	5729813	0.99
2	2107	Iron Man	174	1	3	A. F. Iommi, W. Ward, T. Butler, J. Osbourne	172120	5609799	0.99
3	1908	New Rhumba	157	1	2	A. Jamal	276871	8980400	0.99
4	415	Astronomy	35	1	3	A. Bouchard/J. Bouchard/S. Pearlman	397531	13065612	0.99
5	2589	Hard To Handle	210	1	6	A. Isbell/A. Jones/O. Redding	206994	6786304	0.99
6	3427	Fanfare for the Common Man	296	2	24	Aaron Copland	198064	3211245	0.99
7	3357	OAM's Blues	267	5	2	Aaron Goldberg	266936	4292028	0.99
8	20	Overdose	4	1	1	AC/DC	369319	12066294	0.99
9	17	Let There Be Rock	4	1	1	AC/DC	366654	12021261	0.99
10	15	Go Down	4	1	1	AC/DC	331180	10847611	0.99
11	19	Problem Child	4	1	1	AC/DC	325041	10617116	0.99



Question Action

Sort invoices by billing_city (ascending) and total purchase (descending), show the invoice_id, billing_city and total

```
In [77]: pd.read_sql_query("""
    SELECT invoice_id, billing_city, total
    FROM invoice
    ORDER BY billing_city ASC, total DESC
""", engine)
```

Out[77]:

	invoice_id	billing_city	total
0	390	Amsterdam	13.86
1	206	Amsterdam	8.94
2	32	Amsterdam	8.91
3	184	Amsterdam	3.96
4	379	Amsterdam	1.98
...
407	388	Yellowknife	5.94
408	366	Yellowknife	3.96
409	343	Yellowknife	1.98
410	148	Yellowknife	1.98
411	27	Yellowknife	0.99

412 rows × 3 columns

Distinct

Selects only the unique values of a variable

```
In [79]: # Look at the Distinct cities in our customer list
```

```
pd.read_sql_query("""SELECT DISTINCT city
                   FROM customer
                   ORDER BY city
                   LIMIT 20;"""
                   ,engine)
```

Out[79]:

	city
0	Amsterdam
1	Bangalore
2	Berlin
3	Bordeaux
4	Boston
5	Brasília
6	Brussels
7	Budapest
8	Buenos Aires
9	Chicago
10	Copenhagen
11	Cupertino
12	Delhi
13	Dijon
14	Dublin
15	Edinburgh
16	Edmonton
17	Fort Worth
18	Frankfurt
19	Halifax

Question/Action

Find the list of distinct artists listed in Track, sort them

```
In [81]: pd.read_sql_query("""
    SELECT DISTINCT Composer
    FROM Track
    ORDER BY Composer
""", engine)
```

Out[81]:

	composer
0	A. F. Iommi, W. Ward, T. Butler, J. Osbourne
1	A. Jamal
2	A.Bouchard/J.Bouchard/S.Pearlman
3	A.Isbell/A.Jones/O.Redding
4	Aaron Copland
...	...
849	Willie Dixon
850	Willie Dixon, C. Burnett
851	Wolfgang Amadeus Mozart
852	Wright, Waters
853	None

854 rows × 1 columns

Where

Where is a filter that allows us to filter out only the rows that meet some desired condition.

Notice that the select command itself allows us to control the columns show, Where works on the rows

Comparison Operators

=, !=, <, >, >=, <= *Note equality is a single equal sign in postgres "="

Logical Operators

AND, NOT, OR

Other tests

ALL- 1 if all expressions are 1

ANY- 1 if any expressions is 1

BETWEEN- tests for a range of values

IN- comparison to a list of values

LIKE- used on strings, if they match a pattern

```
In [83]: # we can select a specific album id for the tracks  
  
pd.read_sql_query("""SELECT name, milliseconds, bytes, album_id  
                  FROM track  
                  WHERE album_id=6""", engine)
```

Out[83]:

	name	milliseconds	bytes	album_id
0	All I Really Want	284891	9375567	6
1	You Oughta Know	249234	8196916	6
2	Perfect	188133	6145404	6
3	Hand In My Pocket	221570	7224246	6
4	Right Through You	176117	5793082	6
5	Forgiven	300355	9753256	6
6	You Learn	239699	7824837	6
7	Head Over Feet	267493	8758008	6
8	Mary Jane	280607	9163588	6
9	Ironic	229825	7598866	6
10	Not The Doctor	227631	7604601	6
11	Wake Up	293485	9703359	6
12	You Oughta Know (Alternate)	491885	16008629	6

In [85]: *# we can select a specific album id for the tracks and restrict to relatively short tracks*

```
pd.read_sql_query("""SELECT name, milliseconds, bytes, album_id  
                  FROM track  
                 WHERE album_id=6 AND milliseconds<250000""", engine)
```

Out[85]:

		name	milliseconds	bytes	album_id
0	You Oughta Know		249234	8196916	6
1	Perfect		188133	6145404	6
2	Hand In My Pocket		221570	7224246	6
3	Right Through You		176117	5793082	6
4	You Learn		239699	7824837	6
5	Ironic		229825	7598866	6
6	Not The Doctor		227631	7604601	6

In []: # *Question/Action*

Find out how many invoices totals where over 25

In [88]:

```
# Query to count invoices with totals over 25
pd.read_sql_query("""
    SELECT COUNT(*) AS "Invoices_Over_25"
    FROM Invoice
    WHERE total > 25
    -- Changed "Total" to total to match the actual column name in the database
""", engine)
```

Out[88]:

Invoices_Over_25

0	1
---	---

LIKE

In []: # *Question/Action*

Use the LIKE function to find all the invoice entries **from** Ireland

```
be sure to use LIKE, the = test would work here too, but practice using LIKE
```

```
In [96]: pd.read_sql_query("SELECT column_name FROM information_schema.columns WHERE table_name='invoice'", engine)
```

```
Out[96]:
```

	column_name
0	invoice_id
1	customer_id
2	invoice_date
3	total
4	billing_city
5	billing_state
6	billing_country
7	billing_postal_code
8	billing_address

```
In [98]: # The Like operator, allows partial text matching
```

```
# note the use of the doubled percent signs %%
# also note that this is case sensitive

pd.read_sql_query("""SELECT customer_id, billing_city, billing_country
                   FROM invoice
                   WHERE billing_country LIKE '%Ireland%'""", engine)
```

Out[98]:

	customer_id	billing_city	billing_country
0	46	Dublin	Ireland
1	46	Dublin	Ireland
2	46	Dublin	Ireland
3	46	Dublin	Ireland
4	46	Dublin	Ireland
5	46	Dublin	Ireland
6	46	Dublin	Ireland

IN

Tests for membership in a list

Also filtering out one AC/DC album using AND NOT combined with LIKE

In [100...]

```
pd.read_sql_query("""SELECT
    name,
    album_id,
    media_type_id
  FROM
    track
  WHERE
    media_type_id IN (2, 3) AND NOT(name LIKE '%%Wall%%');""",engine)
```

Out[100...]

		name	album_id	media_type_id
0		Fast As a Shark	3	2
1		Restless and Wild	3	2
2		Princess of the Dawn	3	2
3		Welcome to the Jungle	90	2
4		It's So Easy	90	2
...	
444		There's No Place Like Home, Pt. 2	261	3
445		There's No Place Like Home, Pt. 3	261	3
446		Band Members Discuss Tracks from "Revelations"	271	3
447		Branch Closing	251	3
448		The Return	251	3

449 rows × 3 columns

AND

In [102...]

```
pd.read_sql_query("""SELECT
    billing_address,
    billing_city,
    total
FROM
    invoice
WHERE
    billing_city= 'New York'
AND total > 5
ORDER BY
    total;""",engine)
```

Out[102...]

	billing_address	billing_city	total
0	627 Broadway	New York	5.94
1	627 Broadway	New York	8.91
2	627 Broadway	New York	13.86

In [50]: `pd.read_sql_query("""SELECT * FROM invoice LIMIT 5""",engine)`

Out[50]:

	invoice_id	customer_id	invoice_date	billing_address	billing_city	billing_state	billing_country	billing_postal_code	total
0	1	2	2021-01-01	Theodor-Heuss-Straße 34	Stuttgart	None	Germany	70174	1.98
1	2	4	2021-01-02	Ullevålsveien 14	Oslo	None	Norway	0171	3.96
2	3	8	2021-01-03	Grétrystraat 63	Brussels	None	Belgium	1000	5.94
3	4	14	2021-01-06	8210 111 ST NW	Edmonton	AB	Canada	T6G 2C7	8.91
4	5	23	2021-01-11	69 Salem Street	Boston	MA	USA	2113	13.86

OR

Using AND and OR together

In [104...]

```
pd.read_sql_query("""
    SELECT
        billing_address,
        billing_city,
        total
    FROM
        invoice
    WHERE
        (billing_city= 'New York' OR billing_city= 'Chicago')
    AND total > 5
```

```
        ORDER BY
            total; """",engine)
```

Out[104...]

	billing_address	billing_city	total
0	162 E Superior Street	Chicago	5.94
1	627 Broadway	New York	5.94
2	162 E Superior Street	Chicago	7.96
3	162 E Superior Street	Chicago	8.91
4	627 Broadway	New York	8.91
5	627 Broadway	New York	13.86
6	162 E Superior Street	Chicago	15.86

BETWEEN

Looks for a range of values

In [106...]

```
pd.read_sql_query("""SELECT
                    invoice_id,
                    billing_address,
                    total
                FROM
                    invoice
                WHERE
                    total BETWEEN 14.91 and 18.86
                ORDER BY
                    total; """",engine)
```

Out[106...]

	invoice_id	billing_address	total
0	193	Berger Straße 10	14.91
1	208	Ullevålsveien 14	15.86
2	103	162 E Superior Street	15.86
3	313	68, Rue Jouvence	16.86
4	306	Klanova 9/506	16.86
5	88	Calle Lira, 198	17.91
6	201	319 N. Frances Street	18.86
7	89	Rotenturmstraße 4, 1010 Innere Stadt	18.86

In [108...]

```
#NOT BETWEEN
#
# excluding a range

pd.read_sql_query("""SELECT
                    invoice_id,
                    billing_address,
                    total
                FROM
                    invoice
                WHERE
                    total NOT BETWEEN 1 and 20
                ORDER BY
                    total; """ ,engine)
```

Out[108]:

	invoice_id	billing_address	total
0	405	541 Del Medio Avenue	0.99
1	13	1600 Amphitheatre Parkway	0.99
2	20	110 Raeburn Pl	0.99
3	27	5112 48 Street	0.99
4	34	Praça Pio X, 119	0.99
5	41	C/ San Bernardo 85	0.99
6	48	796 Dundas Street West	0.99
7	55	Grétrystraat 63	0.99
8	62	3 Chatham Street	0.99
9	69	319 N. Frances Street	0.99
10	76	Ullevålsveien 14	0.99
11	83	9, Place Louis Barthou	0.99
12	90	801 W 4th Street	0.99
13	384	162 E Superior Street	0.99
14	391	1498 rue Bélanger	0.99
15	398	11, Place Bellecour	0.99
16	6	Berger Straße 10	0.99
17	104	Barbarossastraße 19	0.99
18	111	1 Microsoft Way	0.99
19	118	421 Bourke Street	0.99
20	125	Rua da Assunção 53	0.99
21	132	Qe 7 Bloco G	0.99

invoice_id		billing_address	total
22	139	Celsiusg. 9	0.99
23	146	230 Elgin Street	0.99
24	153	Sønder Boulevard 51	0.99
25	160	Via Degli Scipioni, 43	0.99
26	167	2211 W Berry Street	0.99
27	174	Klanova 9/506	0.99
28	181	68, Rue Jouvence	0.99
29	188	120 S Orange Ave	0.99
30	195	Av. Brigadeiro Faria Lima, 2170	0.99
31	209	627 Broadway	0.99
32	216	307 Macacha Güemes	0.99
33	223	Rua dos Campeões Europeus de Viena, 4350	0.99
34	230	8210 111 ST NW	0.99
35	237	202 Hoxton Street	0.99
36	244	194A Chain Lake Drive	0.99
37	251	Rua Dr. Falcão Filho, 155	0.99
38	258	Lijnbaansgracht 120bg	0.99
39	265	1033 N Park Ave	0.99
40	272	Rilská 3174/6	0.99
41	279	Porthaninkatu 9	0.99
42	286	69 Salem Street	0.99
43	293	Theodor-Heuss-Straße 34	0.99

invoice_id		billing_address	total
44	300	8, Rue Hanovre	0.99
45	314	Calle Lira, 198	0.99
46	321	Tauentzienstraße 8	0.99
47	328	700 W Pender Street	0.99
48	335	113 Lupus St	0.99
49	342	696 Osborne Street	0.99
50	349	Av. Paulista, 2022	0.99
51	356	Ordynacka 10	0.99
52	363	302 S 700 E	0.99
53	370	Rotenturmstraße 4, 1010 Innere Stadt	0.99
54	377	Erzsébet krt. 58.	0.99
55	96	Erzsébet krt. 58.	21.86
56	194	3 Chatham Street	21.86
57	299	2211 W Berry Street	23.86
58	404	Rilská 3174/6	25.86

```
In [110]: # shut down the engine to close the connection  
engine.dispose()
```

```
In [ ]:
```